

Adding a Custom CPU or GPU Backend to LM Studio

Overview

This document describes how to add a custom `llama.cpp` backend to LM Studio. It applies to both CPU-only builds (with AVX or AVX1) and GPU builds (Vulkan). The procedure allows LM Studio to run on older CPUs, use a custom Vulkan GPU backend, or, analogously, a CUDA backend (requires modern GPU or bfloat16 patching).

Steps Taken

1. Build `llama.cpp` locally.

- CPU build example:

```
ggml-base.dll  
ggml-cpu.dll  
ggml_llamacpp.dll
```

- GPU/Vulkan build example:

```
ggml-base.dll  
ggml-cpu.dll # optional CPU support  
ggml-vulkan.dll  
ggml_llamacpp.dll  
liblmstudio_bindings_vulkan.node  
llm_engine_vulkan.node
```

2. Locate LM Studio backend directory. Example:

```
C:\Users\Admin\.lmstudio\extensions\backends\llama.cpp-win-x86_64-vulkan-avx2  
-1.48.0
```

3. Duplicate and rename the backend folder.

- Use a descriptive name for your build flags. Example:

```
llama.cpp-win-x86_64-vulkan-avx-1.48.0
```

- This prevents overwriting existing backends.

4. Edit backend metadata files.

- `backend-manifest.json`: update `"required_features"` to match your build (e.g., `"AVX"` instead of `"AVX2"`).
- `display-data.json`: update the display name for clarity:

```
"displayName": "Vulkan llama.cpp (Windows, AVX)"
```

- Optional: For CPU-only builds, change display name to:

```
"displayName": "CPU llama.cpp (Windows, AVX)"
```

5. Copy build outputs to backend folder.

- Example PowerShell snippet (CPU or Vulkan):

```
$backend = "$env:USERPROFILE\.lmstudio\extensions\backends\llama.cpp-win-x86_64-vulkan-avx-1.48.0"
$build = "C:\Users\Admin\source\llama.cpp\build-vulkan\bin"

Copy-Item "$build\ggml-base.dll" -Destination $backend -Force
Copy-Item "$build\ggml-cpu.dll" -Destination $backend -Force
Copy-Item "$build\ggml-vulkan.dll" -Destination $backend -Force
# Optional: only overwrite if needed
# Copy-Item "$build\ggml_llamacpp.dll" -Destination $backend -Force
# Copy-Item "$build\liblmstudio_bindings_vulkan.node" -Destination $backend -Force
# Copy-Item "$build\llm_engine_vulkan.node" -Destination $backend -Force
```

- For CPU-only backends, skip Vulkan files.

6. **Refresh LM Studio.** The new backend will appear with your custom display name and use the locally-built binaries.

7. **Optional: CUDA backend.** Repeat the process with a CUDA-configured `llama.cpp` build. Note: Ensure GPU supports bfloat16 or patch the backend for older cards.

Usage Notes

- **For simplicity, just paste the pre built version to backends**
- General idea: copy DLLs from local build to a uniquely-named LM Studio backend folder, adjust metadata, and refresh.
- CPU builds: AVX, AVX1, or no SIMD support can be configured in CMake.
- GPU/Vulkan builds: requires Vulkan SDK; allows running larger models on GPU.
- CUDA builds: requires modern GPU; may need bfloat16 patching for older hardware.
- This method preserves LM Studio stability while using optimized builds.